


Gradient descent: follow the slope downhill to the best weights



Schematic of the optimiser (Cadence uses AdamW). The real loss surface is ~18000-dimensional; this is a 2-D cartoon.